

Introduction to R

Selahattin Murat Sirin

08/09/2020

Outline

- ▶ R Ecosystem
- ▶ R Basics
- ▶ Data Analysis
- ▶ Graphics
- ▶ R Markdown

BASIC R OPERATIONS

R Ecosystem

R is a free statistical computing software. It can be downloaded from CRAN website for free (<https://cran.r-project.org/>). The latest version is R 4.0.2. If you have difficulties in installing the latest version, you can install older versions from (<https://cran.r-project.org/src/base/R-3/>) - the last is the 3.6.3.

RStudio's IDE (Integrated Development Environment) offers a very versatile user interface. There are other IDEs; however, RStudio is the most popular among R users. It can be installed from (<https://rstudio.com/>)

Most people want to use cloud services to save their work. RStudio can interface with Git(Hub). You can create a free GitHub account from (<https://github.com/>), and save your work in there.

Functions and Packages

R is an object-oriented programming which organizes software design around data, or objects, rather than functions and logic. An object can be defined as a data field that has unique attributes and behavior.

Each function has 6 elements:

- 1) Name and description
- 2) Usage
- 3) Arguments
- 4) Details
- 5) Values
- 6) Dependencies and references

```
## Check the sum function  
sum
```

```
?sum
```

All functions are provided through packages, and it provides significant flexibility. On the other hand, unlike other statistical

Example

An example using `install.packages` function

```
install.packages("dplyr", dependencies = TRUE)
```

Load installed package using `library` or `require` functions

```
library(dplyr)  
require(dplyr)
```

You can also load package using `p_load` function provided in **pacman** package. When sharing your work with others, this might be a better solution.

```
install.packages("pacman")  
library(pacman)  
p_load(dplyr)
```

Use `update.packages` function to update packages.

Session Info

Always check your session before starting any project

```
## Provides information about session  
sessionInfo()
```

```
## R version 3.6.2 (2019-12-12)  
## Platform: x86_64-w64-mingw32/x64 (64-bit)  
## Running under: Windows 10 x64 (build 18363)  
##  
## Matrix products: default  
##  
## locale:  
## [1] LC_COLLATE=English_Canada.1252 LC_CTYPE=English_Can  
## [3] LC_MONETARY=English_Canada.1252 LC_NUMERIC=C  
## [5] LC_TIME=English_Canada.1252  
##  
## attached base packages:  
## [1] stats      graphics  grDevices  utils      datasets  me  
##
```

R objects

```
a <- 5; b <- "a"; c <- TRUE  
A <- c(1:5)  
B <- c("a", "b", "c")
```

```
class(a); class(b); class(c)
```

```
## [1] "numeric"
```

```
## [1] "character"
```

```
## [1] "logical"
```

```
class(A); class(B)
```

```
## [1] "integer"
```

```
## [1] "character"
```


Dates

```
date1 <- as.Date("10/23/2016", format = "%m/%d/%Y")  
date2 <- as.Date("23 October 16", format = "%d %B %y")  
date3 <- as.Date("11/10/2016", format = "%m/%d/%Y")
```

```
date1 - date2
```

```
## Time difference of 0 days
```

```
weekdays(date2)
```

```
## [1] "Sunday"
```

```
as.numeric(date3)
```

```
## [1] 17115
```

Besides base operations, **lubridate** package allows multiple operations with dates.

Mathematical operations

```
5 + 5 # Addition
5 - 5 # Subtraction
5 * 5 # Multiplication
5 / 5 # Division
0 / 0 # Division: Not a number (NaN)
pi # Pi
sqrt(5) # Square root
5^2 # Exponent
5**2 # Exponent
5 %% 2 # Division
5%%2 # Modulo
abs(-5) # Absolute value
log(5) # ln
exp(5) # Exponent
log10(5) # log10
log2(5) # log2
round(5/2, digits = 2) # Rounding
signif(12345.6789, digits = 5) # Rounding
signif(5/2) # The built-in function
```

Mathematical operations

Lag and Lead

```
library(dplyr)
x <- seq(1:10)
lag(x,1)
```

```
##   [1] NA  1  2  3  4  5  6  7  8  9
```

```
lag(x,2)
```

```
##   [1] NA NA  1  2  3  4  5  6  7  8
```

```
lead(x,1)
```

```
##   [1]  2  3  4  5  6  7  8  9 10 NA
```

```
lead(x,2)
```

```
##   [1]  3  4  5  6  7  8  9 10 NA NA
```

Mathematical operations

Cumulative Sum, Cumulative Product, Extremes

```
cumsum(x)
```

```
## [1] 1 3 6 10 15 21 28 36 45 55
```

```
cumprod(x)
```

```
## [1] 1 2 6 24 120 720  
## [10] 3628800
```

```
cummax(x)
```

```
## [1] 1 2 3 4 5 6 7 8 9 10
```

```
cummin(x)
```

```
## [1] 1 1 1 1 1 1 1 1 1 1
```

```
cummean(x)
```

```
## [1] 1.0 1.5 2.0 2.5 3.0 3.5 4.0 4.5 5.0 5.5
```

Mathematical operations

```
x <- 1:20  
first(x)
```

```
## [1] 1
```

```
last(x)
```

```
## [1] 20
```

```
nth(x,9)
```

```
## [1] 9
```

```
range(x)
```

```
## [1] 1 20
```

R Operators

- ▶ `<-` Assign value
- ▶ `=` Assign value
- ▶ `:` Generate regular sequences
- ▶ `c()` Combine Values into a Vector or List
- ▶ `(<)` Smaller
- ▶ `(>)` Greater
- ▶ `<=` Smaller or equal
- ▶ `=` *Greater or equal*
- ▶ `==` Equal
- ▶ `!=` Not equal
- ▶ `Or`
- ▶ `&` And

Matrix operations

- ▶ `matrix()` Creates matrix
- ▶ `rbind()` Bind rows
- ▶ `cbind()` Bind columns
- ▶ `ncol()` Number of columns
- ▶ `nrow()` Number of rows
- ▶ `t()` Transpose of a matrix
- ▶ `det()` Determinant of a matrix
- ▶ `eigen()` Eigen values
- ▶ `solve()` Solves equation
- ▶ `diag()` Construct a diagonal matrix
- ▶ `A%*%B` Dot product
- ▶ `A%o%B` Outer product
- ▶ `crossprod(A,B)` Return a matrix cross-product
- ▶ `kronecker()` Computes the generalised kronecker product

Matrix examples

```
A <- matrix(c(2,5,7,1,6,12,3,2,2), byrow = TRUE, ncol = 3)
a <- c(2,5,7)
b <- c(1,6,12)
c <- c(3,2,2)
B <- rbind(a,b,c)
t(A)
solve(A)
eigen(A)
diag(4)
```

To select an element in a matrix you can use `A[2,3]`. This gives the element in the second row and third column

Data structure

- ▶ Vector (Atomic) is the fundamental data structure. There are six basic vector types: (1) Logical, (2) Integer, (3) Double, (4) Character, (5) Complex, (6) Raw. Use `c()` or `vector()` functions to create vectors. The elements of a vector are all of the same type.
- ▶ List is referred as generic vector. The elements may be different types. Use `list()` to create a list.
- ▶ Matrix is a two dimensional array. Use `matrix()` to create a matrix.
- ▶ Array is a multi-dimensional vector. All elements should be the same type. Use `array()` to create an array

Data frame

- Data frame is stored as a list of vectors, factors and/or matrices each of which has the same length. Use `data.frame()` to create a data frame.

```
df <- data.frame(a = c(1:10), b = rnorm(10, 0, 0.5), c = 1:10)
```

Table 1: Data Frame

	a	b	c
1	-0.6170387	1	1
2	-0.2865482	2	2
3	0.4057937	1	3
4	0.6413598	2	4
5	0.1358248	1	5
6	0.0568945	2	6

Basic data functions

- ▶ `head()` returns the first part of a vector
- ▶ `tail()` returns the first part of a vector
- ▶ `names()`, `colnames()` displays column names
- ▶ `dim()` retrieves the dimension of an object
- ▶ `str()` displays the structure of an R Object
- ▶ `table()` used for cross tabulation
- ▶ `xtabs()` creates a contingency table
- ▶ `margin.table()` creates contingency table in array form
- ▶ `prop.table()` express table ntries as a fraction of marginal table
- ▶ `ftable()` creates 'flat' contingency tables

Basic data functions - Examples

```
library(datasets)
data(mtcars)
head(mtcars)
```

##	mpg	cyl	disp	hp	drat	wt	qsec	vs
## Mazda RX4	21.0	6	160	110	3.90	2.620	16.46	0
## Mazda RX4 Wag	21.0	6	160	110	3.90	2.875	17.02	0
## Datsun 710	22.8	4	108	93	3.85	2.320	18.61	1
## Hornet 4 Drive	21.4	6	258	110	3.08	3.215	19.44	1
## Hornet Sportabout	18.7	8	360	175	3.15	3.440	17.02	0
## Valiant	18.1	6	225	105	2.76	3.460	20.22	1

Basic data functions - Examples

In basic R operations, you can use `with()` function to evaluate an expression in an environment constructed from data.

```
## Create cross-table  
with(mtcars, table(cyl, gear))
```

```
##      gear  
## cyl   3   4   5  
##    4   1   8   2  
##    6   2   4   1  
##    8  12   0   2
```

Missing data

- ▶ `complete.cases()` returns a logical vector indicating which cases are complete
- ▶ `na.omit` removes rows with missing data

```
missing.values <- c(1,2,NA, 4, 5, NA, 7, NA, NA)  
complete.cases(missing.values)
```

```
## [1] TRUE TRUE FALSE TRUE TRUE FALSE TRUE FALSE FALSE
```

```
na.omit(missing.values)
```

```
## [1] 1 2 4 5 7
```

```
## attr(,"na.action")
```

```
## [1] 3 6 8 9
```

```
## attr(,"class")
```

```
## [1] "omit"
```

Many packages provide functions for handling missing data. Check **dplyr**, **Amelia**, **Hmisc**, **missForest** packages for examples.

Import Data

- ▶ `read.table(file)`
- ▶ `read.csv(file)`
- ▶ `read.csv2(file)`
- ▶ `read.delim(file)`
- ▶ `read.delim2(file)`

Check **readr**, **XLConnect** packages to import data from excel files

Import Data

You can import data from other statistical programs using **foreign** and **Hmisc** packages

```
library(foreign)
library(Hmisc)
# SPSS
data.spss <- read.spss("File name")

# SAS
data.sas <- sasxport.get("File name")

# Stata
data.stata <- read.stata("File name")

# Systat
data.systat <- read.systat("File name")
```


BASIC DATA OPERATIONS

Rename variable

Renaming a variable using `names()` function.

```
library(datasets)
data(mtcars)
names(mtcars)
```

```
## [1] "mpg" "cyl" "disp" "hp" "drat" "wt" "qsec" "v"
## [11] "carb"
```

```
names(mtcars)[2] = "lyc"
names(mtcars)
```

```
## [1] "mpg" "lyc" "disp" "hp" "drat" "wt" "qsec" "v"
## [11] "carb"
```

Rename variables

Tidyverse is a collection of packages that are essential for data analysis. One of the most useful packages is the **dplyr** package. `rename()` package helps to rename variables.

`%>%` function is the pipeline function called from **magrittr** package. It simplifies analysis to a great extent.

```
library(datasets)
library(dplyr)
data(mtcars)
mtcars <- mtcars %>%
  rename(MilesPGallon = mpg)
names(mtcars)
```

```
##   [1] "MilesPGallon" "cyl"           "disp"          "hp"
##   [6] "wt"           "qsec"          "vs"            "am"
##  [11] "carb"
```

Create a new variable

The base R function has \$ operator that acts on vectors, matrices, arrays and lists to extract or replace parts.

```
library(datasets)
data(mtcars)
mtcars$mpgc <- mtcars$mpg / mtcars$cyl
summary(mtcars$mpgc)
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##   1.300   1.928   3.108   3.837   5.700   8.475
```

We can do the same with mutate() function in **dplyr** package.

```
library(dplyr)
library(datasets)
data(mtcars)
mtcars <- mtcars %>%
  mutate( mpgc = mpg /cyl)
```

Ordering variables

The base `order()` function rearranges the data set. `attach()` function attaches the dataset to the environment, so objects in the database can be accessed by simply giving their names. **dplyr** package has `arrange()` function for the same purpose, and it is more convenient.

```
library(datasets)
data(mtcars)
attach(mtcars)
mtcars[order(mpg, decreasing = TRUE),] %>%
  head()
```

##		mpg	cyl	disp	hp	drat	wt	qsec	vs	am
##	Toyota Corolla	33.9	4	71.1	65	4.22	1.835	19.90	1	1
##	Fiat 128	32.4	4	78.7	66	4.08	2.200	19.47	1	1
##	Honda Civic	30.4	4	75.7	52	4.93	1.615	18.52	1	1
##	Lotus Europa	30.4	4	95.1	113	3.77	1.513	16.90	1	1
##	Fiat X1-9	27.3	4	79.0	66	4.08	1.935	18.90	1	1
##	Porsche 914-2	26.0	4	120.3	91	4.43	2.140	16.70	0	1

Merging datasets

The base `merge()` function merges data sets by common columns or row names. Similarly, `_join()` functions in `dplyr` packages offer more options. Compare these two data sets using `merge()` function.

```
## Create two data frames
```

```
df_1 <- data.frame(id = c(1,2,3), value = rnorm(3,1,0.2))  
df_2 <- data.frame(id = c(3,4,5,6), value = rnorm(4,1,0.2))  
new_df <- merge(df_1, df_2, by = "id")  
new_df_1 <- merge(df_1, df_2, by = "id", all = TRUE)  
print(new_df)
```

```
##   id  value.x  value.y  
## 1   3 0.8031866 1.218372
```

Merging datasets using join functions

```
inner_join(x, y, by = NULL)
left_join(x, y, by = NULL)
right_join(x, y, by = NULL)
full_join(x, y, by = NULL)
anti_join(x, y, by = NULL)
```

```
library(dplyr)
inner_join(df_1, df_2, by = "id")
```

```
##   id   value.x value.y
## 1  3 0.8031866 1.218372
```

Subsetting data

- ▶ `X[n]` select nth element; nth row if `dim > 2`
- ▶ `X[n,k]` select element from nth row, kth column
- ▶ `X[-n]` All but nth element
- ▶ `X[1:n]` First n elements
- ▶ `X[c(a,b,c)]` select multiple elements
- ▶ `X[x > k]` selects elements greater than k
- ▶ `X[x > k | x < k]` selects elements greater than or less

```
library(datasets)
data(mtcars)
mtcars$mpg[10:15]
```

```
## [1] 19.2 17.8 16.4 17.3 15.2 10.4
```

```
mtcars$mpg[mpg > 30]
```

```
## [1] 32.4 30.4 33.9 30.4
```


Subsetting data

The base `subset()` function and `filter()` or `select()` functions from **dplyr** package offers more flexibility when subsetting data.

```
library(datasets)
data(airquality)
subset(airquality, Temp > 80, select = c(Ozone, Temp)) %>%
  head()
```

##	Ozone	Temp
## 29	45	81
## 35	NA	84
## 36	NA	85
## 38	29	82
## 39	NA	87
## 40	71	90

Summarizing data sets

Use `summary()` function to summarize variables

```
data(mtcars)
summary(mtcars[,1:4])
```

##	mpg	cyl	disp	h
##	Min. :10.40	Min. :4.000	Min. : 71.1	Min.
##	1st Qu.:15.43	1st Qu.:4.000	1st Qu.:120.8	1st Qu.
##	Median :19.20	Median :6.000	Median :196.3	Median
##	Mean :20.09	Mean :6.188	Mean :230.7	Mean
##	3rd Qu.:22.80	3rd Qu.:8.000	3rd Qu.:326.0	3rd Qu.
##	Max. :33.90	Max. :8.000	Max. :472.0	Max.

Summarizing data sets

Row and column sums and means can be retrieved by these functions:

```
colSums(x, na.rm = FALSE, dims = 1)
rowSums(x, na.rm = FALSE, dims = 1)
colMeans(x, na.rm = FALSE, dims = 1)
rowMeans(x, na.rm = FALSE, dims = 1)
```

```
library(datasets)
data(mtcars)
colMeans(mtcars)
```

##	mpg	cyl	disp	hp	drat
##	20.090625	6.187500	230.721875	146.687500	3.596563
##	vs	am	gear	carb	
##	0.437500	0.406250	3.687500	2.812500	

Note: There are many packages that offer summary tables such as **stargazer**.

Stargazer package

```
library(stargazer)
library(dplyr)
library(datasets)
data(mtcars)
mtcars %>%
  select(hp, mpg, disp) %>%
  stargazer(., median = TRUE,
            iqr = TRUE, type = "text")
```

##

=====

Statistic N Mean St. Dev. Min Pctl(25) Median Pctl(75)

hp 32 146.688 68.563 52 96.5 123 180

mpg 32 20.091 6.027 10 15.4 19.2 22.8

disp 32 230.722 123.939 71 120.8 196.3 326

Reshaping data sets

To reshape data, **tidyr** package has `pivot_()` functions (`pivot_longer` and `pivot_wider`). Assume we have 3 companies with daily average stock prices in USD.

```
require(tidyr)
set.seed(23)
stocks <- data.frame(
  DATE = as.Date('2017-01-01') + 0:9,
  A = as.integer(rnorm(10, 5, 1)),
  B = as.integer(rnorm(10, 10, 1)),
  C = as.integer(rnorm(10, 15, 1)))
stocks.long <- pivot_longer(stocks, A:C, names_to = "Companies",
                           values_to = "Price")
head(stocks.long)
```

```
## # A tibble: 6 x 3
##   DATE          Companies Price
##   <date>        <chr>    <int>
## 1 2017-01-01 A          5
```

Strings

stringr package offers `str_()` functions to handle strings in R.

```
library(stringr)
str_to_upper(string, locale = "")
str_to_lower(string, locale = "")
str_to_title(string, locale = "")
str_count(string, pattern = "")
str_length(string)
str_detect(string, pattern)
str_extract(string, pattern)
str_replace(string, pattern, replacement)
str_sub(string, start = 1L, end = -1L)
str_trim(string, side = c("both", "left", "right"))
```